

Komal Kumari

Project Name: Prediction for Credit Card Approval

Welcome to the Machine Learning Internship program, focused on Credit Card Approval Prediction. In this project, your task is to utilize machine learning techniques to predict whether an applicant will be approved for a credit card or not.

Problem Statement:

The primary objective of this project is to predict the approval or rejection of credit card applications. The challenge lies in understanding the key factors influencing credit card approval decisions and building a predictive model to assist in the decision-making process.

Your Mission:

1. Exploratory Data Analysis (EDA):
2. Feature Engineering:
3. Data Preprocessing:
4. Machine Learning Model Development:
5. Model Evaluation:
6. ML Pipelines and Model Deployment:
7. Recommendations:

```
In [9]: #IMPORT LIBRARIES
```

```
In [10]: # Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
```

```
In [4]: df = pd.read_csv('train_data.csv')
df
```

Out[4]:

[illegible]

In [11]:

df.head()

Out[11]:

	ID	Gender	Has a car	Has a property	Children count	Income	Employment status	Education level	Marital status	Dwelling	Age	Employment length	Has a mobile phone	Has a work phone	Has a phone	Has an email	
0	5037048	M	Y	Y	0	135000.0	Working	Secondary / secondary special	Married	With parents	-16271	-3111	1	0	0	0	Cr
1	5044630	F	Y	N	1	135000.0	Commercial associate	Higher education	Single / not married	House / apartment	-10130	-1651	1	0	0	0	Acco
2	5079079	F	N	Y	2	180000.0	Commercial associate	Secondary / secondary special	Married	House / apartment	-12821	-5657	1	0	0	0	L
3	5112872	F	Y	Y	0	360000.0	Commercial associate	Higher education	Single / not married	House / apartment	-20929	-2046	1	0	0	1	Ma
4	5105858	F	N	N	0	270000.0	Working	Secondary / secondary special	Separated	House / apartment	-16207	-515	1	0	1	0	

In [13]: df.shape

Out[13]: (29165, 20)

In [14]: df.tail()

Out[14]:

	ID	Gender	Has a car	Has a property	Children count	Income	Employment status	Education level	Marital status	Dwelling	Age	Employment length	Has a mobile phone	Has a work phone	Has a phone	Has an email
29160	5067139	F	N	Y	0	112500.0	Pensioner	Secondary / secondary special	Single / not married	House / apartment	-23400	365243	1	0	1	1
29161	5029193	F	N	Y	1	135000.0	Commercial associate	Secondary / secondary special	Married	House / apartment	-15532	-8256	1	0	0	0
29162	5047710	F	N	Y	0	76500.0	Working	Secondary / secondary special	Married	House / apartment	-17782	-3291	1	1	1	0
29163	5009886	F	N	Y	0	157500.0	Pensioner	Secondary / secondary special	Civil marriage	House / apartment	-21635	365243	1	0	1	0
29164	5062632	F	N	Y	0	585000.0	Commercial associate	Secondary / secondary	Married	House / apartment	-18858	-2010	1	0	1	0

```
df.head()
```

0	5037048	M	Y	Y	0	135000.0	Working	Secondary / secondary special	Married	With parents	-16271	-3111	1	0	0	0	0	Ci
1	5044630	F	Y	N	1	135000.0	Commercial associate	Higher education	Single / not married	House / apartment	-10130	-1651	1	0	0	0	0	Acco
2	5079079	F	N	Y	2	180000.0	Commercial associate	Secondary / secondary special	Married	House / apartment	-12821	-5657	1	0	0	0	0	L
3	5112872	F	Y	Y	0	360000.0	Commercial associate	Higher education	Single / not married	House / apartment	-20929	-2046	1	0	0	1	1	Mi
4	5105858	F	N	N	0	270000.0	Working	Secondary / secondary special	Separated	House / apartment	-16207	-515	1	0	1	0	0	

```
df.shape
```

(29165, 20)

```
df.tail()
```

29160	5067139	F	N	Y	0	112500.0	Pensioner	Secondary / secondary special	Single / not married	House / apartment	-23400	365243	1	0	1	1
29161	5029193	F	N	Y	1	135000.0	Commercial associate	Secondary / secondary special	Married	House / apartment	-15532	-8256	1	0	0	0
Secondary																

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29165 entries, 0 to 29164
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     29165 non-null  int64
1   Gender                               29165 non-null  object
2   Has a car                             29165 non-null  object
3   Has a property                         29165 non-null  object
4   Children count                        29165 non-null  int64
5   Income                                29165 non-null  float64
6   Employment status                     29165 non-null  object
7   Education level                       29165 non-null  object
8   Marital status                        29165 non-null  object
9   Dwelling                              29165 non-null  object
10  Age                                   29165 non-null  int64
11  Employment length                     29165 non-null  int64
12  Has a mobile phone                    29165 non-null  int64
13  Has a work phone                      29165 non-null  int64
14  Has a phone                           29165 non-null  int64
15  Has an email                          29165 non-null  int64
16  Job title                             20138 non-null  object
17  Family member count                   29165 non-null  float64
18  Account age                           29165 non-null  float64
19  Is high risk                          29165 non-null  int64
dtypes: float64(3), int64(9), object(8)
memory usage: 4.5+ MB
```

[illegible]

```
In [19]: df.describe()
```

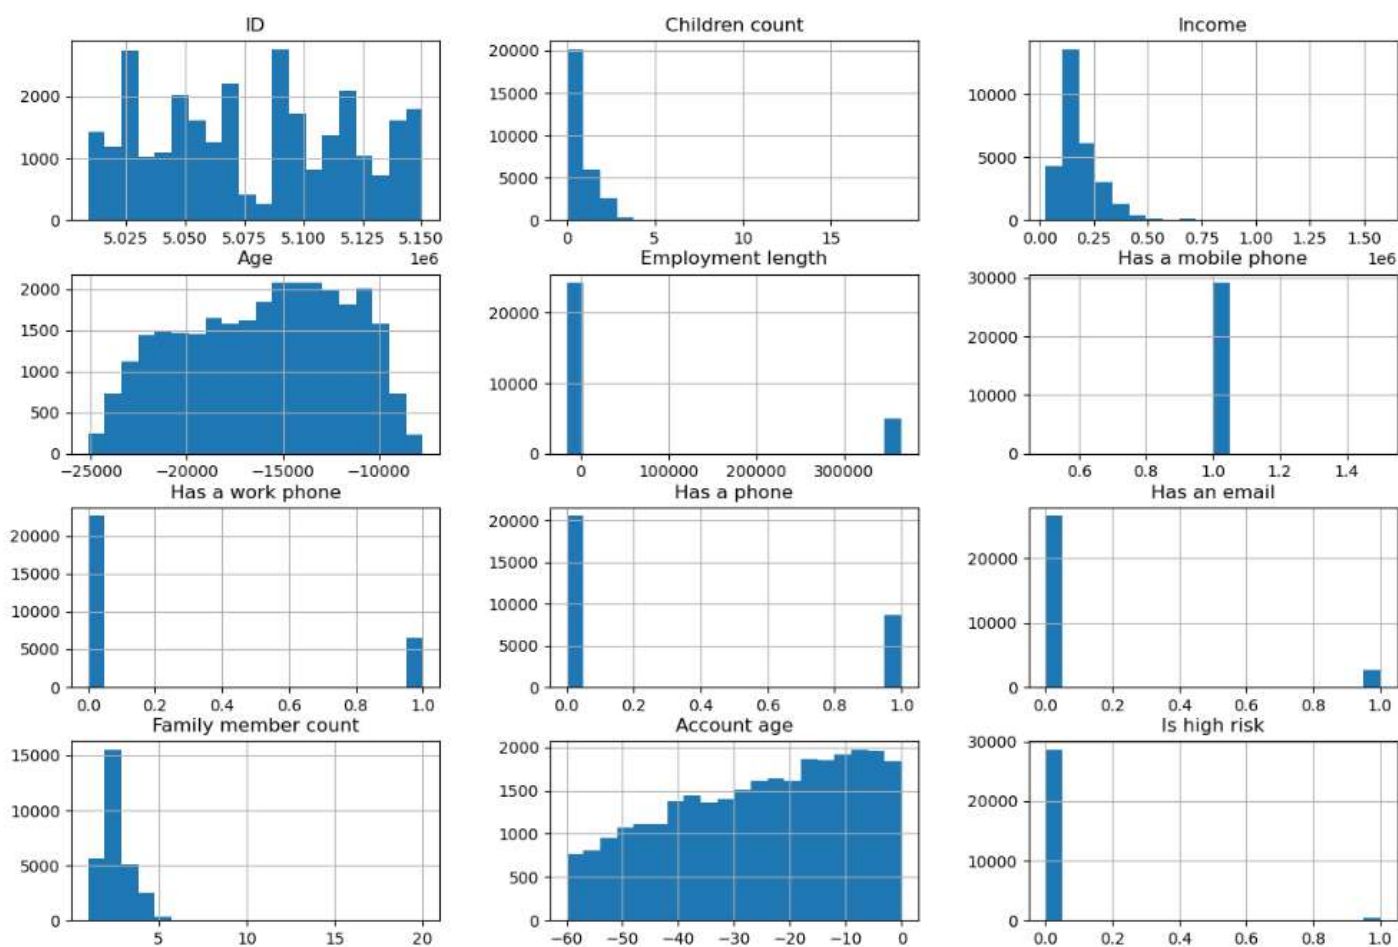
Out[19]:

	ID	Children count	Income	Age	Employment length	Has a mobile phone	Has a work phone	Has a phone	Has an email	Family member count	Account age
count	2.916500e+04	29165.000000	2.916500e+04	29165.000000	29165.000000	29165.0	29165.000000	29165.000000	29165.000000	29165.000000	29165.000000
mean	5.078232e+06	0.430790	1.868904e+05	-15979.477490	59257.761255	1.0	0.224310	0.294977	0.090279	2.197531	-26.137734
std	4.182400e+04	0.741882	1.014096e+05	4202.997485	137655.883458	0.0	0.417134	0.456040	0.286587	0.912189	16.486702
min	5.008804e+06	0.000000	2.700000e+04	-25152.000000	-15713.000000	1.0	0.000000	0.000000	0.000000	1.000000	-60.000000
25%	5.042047e+06	0.000000	1.215000e+05	-19444.000000	-3153.000000	1.0	0.000000	0.000000	0.000000	2.000000	-39.000000
50%	5.074666e+06	0.000000	1.575000e+05	-15565.000000	-1557.000000	1.0	0.000000	0.000000	0.000000	2.000000	-24.000000
75%	5.114629e+06	1.000000	2.250000e+05	-12475.000000	-412.000000	1.0	0.000000	1.000000	0.000000	3.000000	-12.000000
max	5.150485e+06	19.000000	1.575000e+06	-7705.000000	365243.000000	1.0	1.000000	1.000000	1.000000	20.000000	0.000000

```
In [20]: # Checking for missing values
print(df.isnull().sum())
```

```
ID          0
Gender       0
Has a car    0
Has a property 0
Children count 0
Income       0
Employment status 0
Education level 0
Marital status 0
Dwelling     0
Age          0
Employment length 0
Has a mobile phone 0
Has a work phone 0
Has a phone  0
Has an email  0
Job title    9027
Family member count 0
Account age  0
Is high risk  0
dtype: int64
```

```
In [21]: # Visualize the distribution of numerical features
df.hist(bins=20, figsize=(15,10))
plt.show()
```




```
In [68]: # Count plot for categorical features
plt.figure(figsize =(5,5))
for col in ['Gender', 'Has a car', 'Has a property', 'Employment status', 'Education level', 'Marital status', 'Dwelling']:
    sns.countplot(y=col, data=df)
    plt.show()
```



4. Feature Engineering

```
In [23]: # Creating interaction terms (example)
df['Income_per_family_member'] = df['Income'] / (df['Family member count'] + 1) # Avoid division by zero

# Creating binary features (example)
df['Is_Employed'] = df['Employment status'].apply(lambda x: 1 if x != 'Unemployed' else 0)
```

```
In [ ]:
```

5. Data Preprocessing

```
In [81]: df.isnull().sum()
```

5. Data Preprocessing

```
In [81]: df.isnull().sum()
```

```
Out[81]: ID                0
Gender                0
Has a car              0
Has a property         0
Children count         0
Income                0
Employment status      0
Education level        0
Marital status         0
Dwelling               0
Age                   0
Employment length      0
Has a mobile phone     0
Has a work phone       0
Has a phone            0
Has an email           0
Job title              9027
Family member count    0
Account age            0
Is high risk           0
Income_per_family_member 0
Is_Employed            0
dtype: int64
```

```
In [27]: duplicates = df['Job title'].duplicated(keep=False)
print(df[duplicates]) # Print rows where 'Job title' is duplicated
```

	ID	Gender	Has a car	Has a property	Children count	Income	\
0	5037048	M	Y	Y	0	135000.0	
1	5044630	F	Y	N	1	135000.0	
2	5079079	F	N	Y	2	180000.0	
3	5112872	F	Y	Y	0	360000.0	
4	5105858	F	N	N	0	270000.0	
...	
29160	5067139	F	N	Y	0	112500.0	
29161	5029193	F	N	Y	1	135000.0	
29162	5047710	F	N	Y	0	76500.0	
29163	5009886	F	N	Y	0	157500.0	
29164	5062632	F	N	Y	0	585000.0	

29160	Single / not married	House / apartment	...	1
29161	Married	House / apartment	...	1
29162	Married	House / apartment	...	1
29163	Civil marriage	House / apartment	...	1
29164	Married	House / apartment	...	1

	Has a work phone	Has a phone	Has an email	Job title \
0	0	0	0	Core staff
1	0	0	0	Accountants
2	0	0	0	Laborers
3	0	0	1	Managers
4	0	1	0	NaN
...
29160	0	1	1	NaN
29161	0	0	0	Core staff
29162	1	1	0	Managers
29163	0	1	0	NaN
29164	0	1	0	NaN

	Family member count	Account age	Is high risk \
0	2.0	-17.0	0
1	2.0	-1.0	0
2	4.0	-38.0	0
3	1.0	-11.0	0
4	1.0	-41.0	0
...
29160	1.0	-5.0	0
29161	3.0	-24.0	0
29162	2.0	-29.0	0
29163	2.0	-37.0	0
29164	2.0	-43.0	0

	Income_per_family_member	Is_Employed
0	45000.0	1
1	45000.0	1
2	36000.0	1
3	180000.0	1
4	135000.0	1
...
29160	56250.0	1
29161	33750.0	1
29162	25500.0	1
29163	52500.0	1
29164	195000.0	1

[29165 rows x 22 columns]

```
In [28]: mode_job_title = df['Job title'].mode()[0]
```

```
In [29]: df.loc[duplicates, 'Job title'] = mode_job_title
```

```
In [30]: df.isnull().sum()
```


```
Out[30]: ID                0
Gender                0
Has a car             0
Has a property        0
Children count        0
Income               0
Employment status     0
Education level       0
Marital status        0
Dwelling              0
Age                  0
Employment length     0
Has a mobile phone    0
Has a work phone      0
Has a phone           0
Has an email          0
Job title             0
Family member count   0
Account age           0
Is high risk          0
Income_per_family_member 0
Is_Employed           0
dtype: int64
```

```
In [37]: # Grabbing all the columns from the dataset
col_list = df.columns

# we need only the numerical data
for x in col_list:
    if df[x].dtype != "object":
        sns.boxplot(df[x])
        plt.xlabel(x)
        plt.ylabel("count")
        plt.show()
```

```
In [37]: # Grabbing all the columns from the dataset
col_list = df.columns
```

```
# we need only the numerical data
for x in col_list:
    if df[x].dtype != "object":
        sns.boxplot(df[x])
        plt.xlabel(x)
        plt.ylabel("count")
        plt.show()
```



```
In [38]: # Label Encoding for categorical variables
# Label Encoding for categorical variables
from sklearn.preprocessing import LabelEncoder

# List of columns to be label encoded
label_encode_cols = ['Gender', 'Has a car', 'Has a property', 'Employment status',
                     'Education level', 'Marital status', 'Dwelling', 'Job title']

label_encoders = {} # Dictionary to store label encoders for each column

for col in label_encode_cols:
    label_encoders[col] = LabelEncoder()
    df[col] = label_encoders[col].fit_transform(df[col])

df.head() # Check the first few rows to ensure the encoding worked
```

t[38]:

	ID	Gender	Has a car	Has a property	Children count	Income	Employment status	Education level	Marital status	Dwelling	...	Has a mobile phone	Has a work phone	Has a phone	Has an email	Job title	Family member count	Account age	h
0	5037048	1	1	1	0	135000.0	4	4	1	5	...	1	0	0	0	0	2.0	-17.0	
1	5044630	0	1	0	1	135000.0	0	1	3	1	...	1	0	0	0	0	2.0	-1.0	
2	5079079	0	0	1	2	180000.0	0	4	1	1	...	1	0	0	0	0	4.0	-38.0	
3	5112872	0	1	1	0	360000.0	0	1	3	1	...	1	0	0	1	0	1.0	-11.0	
4	5105858	0	0	0	0	270000.0	4	4	2	1	...	1	0	1	0	0	1.0	-41.0	

5 rows × 22 columns

[39]:

```
# Ensure binary columns are integers (0 and 1)
binary_cols = ['Has a mobile phone', 'Has a work phone', 'Has a phone', 'Has an email']

for col in binary_cols:
    df[col] = df[col].astype(int)
```

6. Split the Dataset

[40]:

```
# Split the dataset into features and target variable
X = df.drop(['ID', 'Is high risk'], axis=1) # Drop target variable and any ID columns
y = df['Is high risk']

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

7. Train Machine Learning Models

[41]:

```
# Initialize models
log_reg = LogisticRegression()
dec_tree = DecisionTreeClassifier()
```

7. Train Machine Learning Models

```
In [41]: # Initialize models
log_reg = LogisticRegression()
dec_tree = DecisionTreeClassifier()
rand_forest = RandomForestClassifier()
grad_boost = GradientBoostingClassifier()
```

```
In [42]: # Train the models
log_reg.fit(X_train, y_train)
dec_tree.fit(X_train, y_train)
rand_forest.fit(X_train, y_train)
grad_boost.fit(X_train, y_train)
```

Out[42]: GradientBoostingClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [43]: # Function to evaluate models
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred))
    print("Classification Report:")
    print(classification_report(y_test, y_pred))
    print(f"ROC AUC Score: {roc_auc_score(y_test, y_pred):.2f}")
    fpr, tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test)[:,-1])
    plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc_score(y_test, y_pred):.2f})')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve')
    plt.legend()
    plt.show()
```

```
In [44]: # Evaluate each model
print("Logistic Regression")
evaluate_model(log_reg, X_test, y_test)
```


```
In [44]: # Evaluate each model
print("Logistic Regression")
evaluate_model(log_reg, X_test, y_test)
```


Logistic Regression
Accuracy: 0.98
Confusion Matrix:
[[8614 0]
 [136 0]]
Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	8614
1	0.00	0.00	0.00	136
accuracy			0.98	8750
macro avg	0.49	0.50	0.50	8750
weighted avg	0.97	0.98	0.98	8750

ROC AUC Score: 0.50

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))






In [45]:

```
print("Decision Tree")
evaluate_model(dec_tree, X_test, y_test)
```

```
Decision Tree
Accuracy: 0.97
Confusion Matrix:
[[8480 134]
 [ 113  23]]
Classification Report:
              precision    recall  f1-score   support

     0       0.99         0.98         0.99         8614
     1       0.15         0.17         0.16          136

 accuracy          0.97
 macro avg          0.57
```



```
In [95]: print("Random Forest")
evaluate_model(rand_forest, X_test, y_test)
```

```
Random Forest
Accuracy: 0.98
Confusion Matrix:
[[8587  27]
 [ 118  18]]
Classification Report:
              precision    recall  f1-score   support

     0       0.99         1.00         0.99         8614
     1       0.40         0.13         0.20          136

 accuracy          0.98         0.98         0.98         8750
 macro avg         0.69         0.56         0.60         8750
 . . . . .
```

Accuracy: 0.98


Confusion Matrix:

```
[[8587  27]
 [ 118  18]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	8614
1	0.40	0.13	0.20	136
accuracy			0.98	8750
macro avg	0.69	0.56	0.60	8750
weighted avg	0.98	0.98	0.98	8750

ROC AUC Score: 0.56



```
In [10]: # Importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
```

```
In [ ]:
```

```
In [4]: df = pd.read_csv('train_data.csv')
df
```

Out[4]:

	ID	Gender	Has a car	Has a property	Children count	Income	Employment status	Education level	Marital status	Dwelling	Age	Employment length	Has a mobile phone	Has a work phone	Has a phone	Has an email
0	5037048	M	Y	Y	0	135000.0	Working	Secondary / secondary special	Married	With parents	-16271	-3111	1	0	0	0
1	5044630	F	Y	N	1	135000.0	Commercial associate	Higher education	Single / not	House / apartment	-10130	-1651	1	0	0	0